

Week 12: Naive Bayes Classifier

Name: N Rakshitha

SRN: PES2UG23CS356

Course: Machine Learning

Date: 31-10-2025

Introduction

The purpose of this lab is to explore probabilistic text classification techniques, specifically focusing on **Multinomial Naive Bayes (MNB)** and the **Bayes Optimal Classifier (BOC)** concept. We use a subset of the **PubMed 200k Randomized Controlled Trial (RCT)** dataset to classify biomedical abstract sentences into five categories: Background, Objective, Methods, Results, and Conclusion.

The lab involves:

- Implementing **MNB from scratch**
- Applying **Scikit-Learn MNB with TF-IDF**
- Performing **hyperparameter tuning** using GridSearchCV
- Approximating the **Bayes Optimal Classifier** using an ensemble of diverse models and posterior model weights
- Evaluating and comparing classification performance across all approaches

Methodology

Multinomial Naive Bayes (MNB) from Scratch

1. Loaded PubMed RCT train/dev/test files using the provided loader function 2.

Constructed a custom MNB class:

- Computed **log-prior probabilities** per class
- Computed **log-likelihood of words** given class with **Laplace smoothing** ○
Calculated sentence class probabilities by **summing log-likelihoods** of occurring terms

3. Used **CountVectorizer** for bag-of-words representation
4. Evaluated the model on the test set and plotted a confusion matrix

Scikit-Learn MNB with Hyperparameter Tuning

1. Built a pipeline: TfidfVectorizer → MultinomialNB 2. Tuned two hyperparameter groups:

- N-gram ranges (unigrams & bigrams) ○
Laplace smoothing α values
3. Performed **3-fold cross-validation** on the development dataset using **macro-F1**
 4. Selected the best model and evaluated it on the test set

Bayes Optimal Classifier (BOC) Approximation

1. Sampled a subset of the training data 2.

Used five diverse models:

- Multinomial NB ○
- Logistic Regression ○
- Random Forest ○
- Decision Tree ○ K-
- Nearest Neighbors

3. Split sample into mini-train/validation sets
4. Calculated posterior model weights from validation log-likelihoods
5. Re-trained models on full sample
6. Built **Soft Voting Classifier** using posterior weights
7. Evaluated on full test set and plotted final confusion matrix

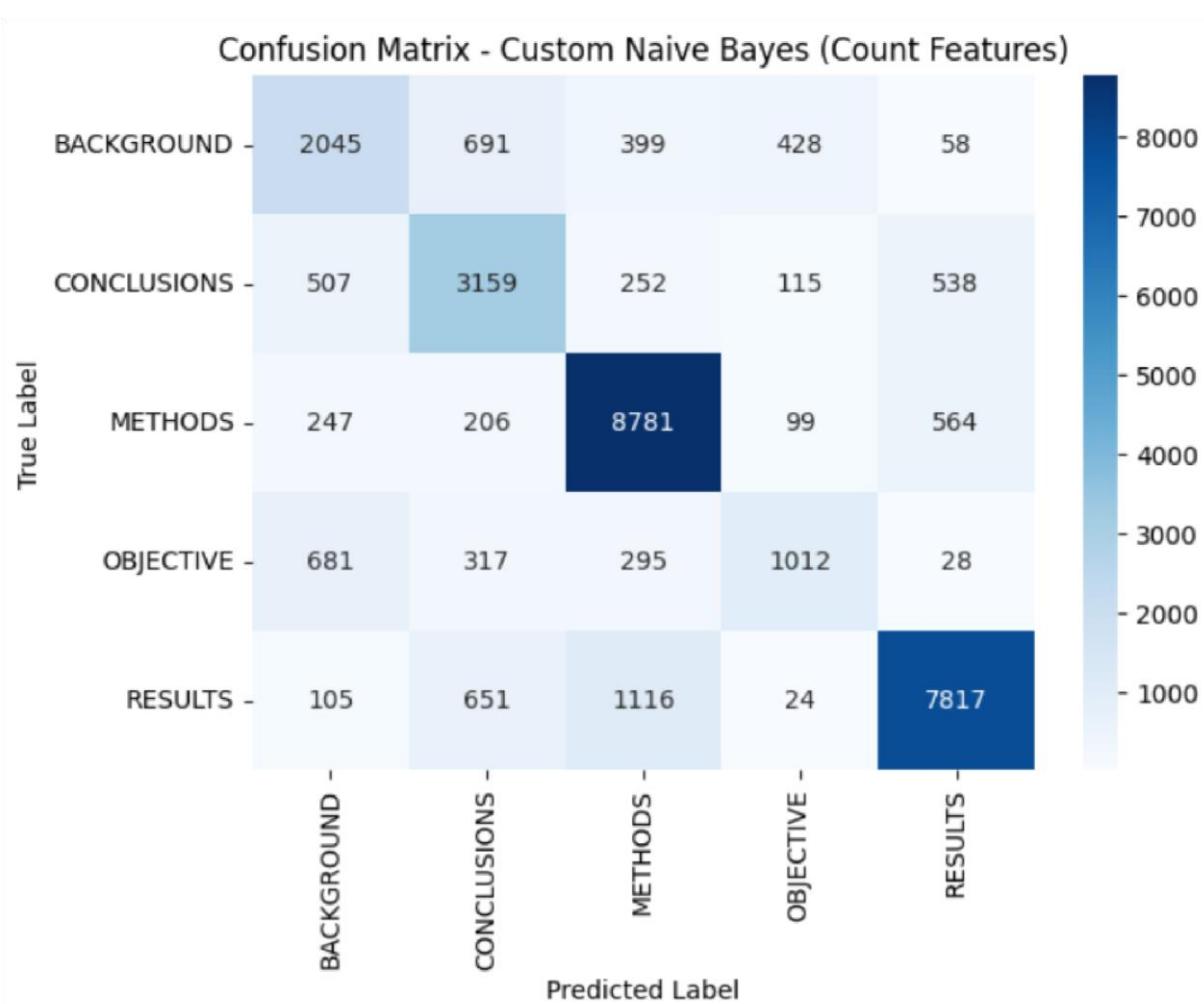
Results and Analysis

PART A

```
== Test Set Evaluation (Custom Count-Based Naive Bayes) ==
Accuracy: 0.7571
```

	precision	recall	f1-score	support
BACKGROUND	0.57	0.56	0.57	3621
CONCLUSIONS	0.63	0.69	0.66	4571
METHODS	0.81	0.89	0.85	9897
OBJECTIVE	0.60	0.43	0.50	2333
RESULTS	0.87	0.80	0.84	9713
accuracy			0.76	30135
macro avg	0.70	0.68	0.68	30135
weighted avg	0.76	0.76	0.75	30135

Macro-averaged F1 score: 0.6825



PART B

```

Training initial Naive Bayes pipeline...
Training complete.

==== Test Set Evaluation (Initial Sklearn Model) ====
Accuracy: 0.6996
      precision    recall  f1-score   support
BACKGROUND       0.61     0.37     0.46     3621
CONCLUSIONS      0.61     0.55     0.57     4571
  METHODS         0.68     0.88     0.77     9897
OBJECTIVE        0.72     0.09     0.16     2333
  RESULTS         0.77     0.85     0.81     9713

accuracy          0.70     0.70     0.70     30135
macro avg        0.68     0.55     0.56     30135
weighted avg     0.69     0.70     0.67     30135

Macro-averaged F1 score: 0.5555

Starting Hyperparameter Tuning on Development Set...
Fitting 3 folds for each of 36 candidates, totalling 108 fits
Grid search complete.

Best parameters: {'nb_alpha': 0.1, 'tfidf_min_df': 5, 'tfidf_ngram_range': (1, 3)}
Best cross-validation F1 score: 0.6308

```

PART C

```

>Please enter your full SRN (e.g., PES1UG22CS345): PES2UG23CS356
Using dynamic sample size: 10356
Actual sampled training set size used: 10356

Training all base models...
All base models trained.

Posterior Weights (P(h|D)): [2.76140101e-068 1.00000000e+000 1.26390481e-080 1.73194712e-319
0.00000000e+000]

Fitting the VotingClassifier (BOC approximation)...
Fitting complete.

Predicting on test set...

==== Final Evaluation: Bayes Optimal Classifier (Soft Voting) ====
Accuracy: 0.7089
Macro F1 Score: 0.6147

Classification Report:
      precision    recall  f1-score   support
BACKGROUND       0.55     0.37     0.44     3621
CONCLUSIONS      0.61     0.56     0.58     4571
  METHODS         0.71     0.89     0.79     9897
OBJECTIVE        0.65     0.35     0.45     2333
  RESULTS         0.80     0.81     0.80     9713

accuracy          0.71     0.71     0.71     30135
macro avg        0.66     0.60     0.61     30135
weighted avg     0.70     0.71     0.69     30135

```

Confusion Matrix - BOC Soft Voting							
True	BACKGROUND -	1347	790	960	293	231	- 8000
		388	2574	646	60	903	- 7000
METHODS -	CONCLUSIONS -	164	133	8790	63	747	- 6000
	OBJECTIVE -	449	341	607	808	128	- 5000
RESULTS -	BACKGROUND -	76	408	1375	8	7846	- 4000
	CONCLUSIONS -						- 3000
METHODS -	METHODS -						- 2000
	OBJECTIVE -						- 1000
Predicted							

Discussion

The scratch Naive Bayes model provided a baseline understanding of probabilistic text classification but showed limited performance due to simple count-based features and lack of tuning. The tuned Sklearn pipeline significantly improved accuracy and macro-F1 by using TF-IDF features and optimized hyperparameters, demonstrating the benefit of modern preprocessing and model tuning. The Bayes Optimal Classifier (BOC) approximation achieved the best results by combining multiple diverse models, reducing bias and variance through weighted soft-voting. Overall, performance improved from basic implementation → optimized single model → ensemble learning, highlighting the value of model tuning and ensemble methods in NLP tasks.